

4. Build a Linear model for the given dataset with regularisation. Attempt to interpret the variable coefficients of the Linear Model

Linear Regression

Importing Libraries

```
In [1]: # Importing Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

Importing the data

```
In [2]: # Importing Data
data = pd.read_csv('nyc_taxi_trip_duration.csv')
data.head()
```

```
Out[2]:
```

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	store_and_fwd_flag	trip_duration
0	id1080784	2	2016-02-29 16:40:21	2016-02-29 16:47:01	1	-73.953918	40.778873	-73.963875	40.771164	N	400
1	id0889885	1	2016-03-11 23:35:37	2016-03-11 23:53:57	2	-73.988312	40.731743	-73.994751	40.694931	N	1100
2	id0857912	2	2016-02-21 17:59:33	2016-02-21 18:26:48	2	-73.997314	40.721458	-73.948029	40.774918	N	1635
3	id3744273	2	2016-01-05 09:44:31	2016-01-05 10:03:32	6	-73.961670	40.759720	-73.956779	40.780628	N	1141
4	id0232939	1	2016-02-17 06:42:23	2016-02-17 06:56:31	1	-74.017120	40.708469	-73.988182	40.740631	N	848

creating a function to identify outliers:

```
In [3]: def UVA_outlier(data, var):
#     import pdb
#     pdb.set_trace()
# calculating descriptives of variable
quant25 = data[var].quantile(0.25)
quant75 = data[var].quantile(0.75)
IQR = quant75 - quant25
med = data[var].median()
whis_low = quant25-(1.5*IQR)
whis_high = quant75+(1.5*IQR)

ls = data.index[(data[var] < whis_low) | (data[var] > whis_high)]

return ls
```

Function to remove outliers

```
In [6]: def remove(df,ls):
ls = sorted(set(ls))
df = df.drop(ls)
return df
```

```
In [16]: # import pdb
index_list1 = []

# for j in data.drop(['id', 'vendor_id', 'pickup_datetime', 'dropoff_datetime', 'store_and_fwd_flag'], axis=1).columns:
for j in ['trip_duration', 'pickup_longitude', 'dropoff_longitude', 'pickup_latitude', 'dropoff_latitude']:
# for j in data.columns:
#     pdb.set_trace()
    for i in [j]:
        index_list1.extend(UVA_outlier(data,i))
        data_cleaned = remove(data,index_list1)
        index_list1.clear()
```

```
In [17]: data_cleaned.shape
```

```
Out[17]: (693076, 11)
```

```
In [18]: data_cleaned.head()
```

```
Out[18]:
```

	id	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	store_and_fwd_flag	trip_duration
0	id1080784	2	2016-02-29 16:40:21	2016-02-29 16:47:01	1	-73.953918	40.778873	-73.963875	40.771164	N	400
1	id0889885	1	2016-03-11 23:35:37	2016-03-11 23:53:57	2	-73.988312	40.731743	-73.994751	40.694931	N	1100
2	id0857912	2	2016-02-21 17:59:33	2016-02-21 18:26:48	2	-73.997314	40.721458	-73.948029	40.774918	N	1635
3	id3744273	2	2016-01-05 09:44:31	2016-01-05 10:03:32	6	-73.961670	40.759720	-73.956779	40.780628	N	1141
4	id0232939	1	2016-02-17 06:42:23	2016-02-17 06:56:31	1	-74.017120	40.708469	-73.988182	40.740631	N	848

```
In [19]: data = data_cleaned
```

Extracting data from date time column

```
In [20]: # # creating an instance(date) of DatetimeIndex class using "pickup_datetime"
date_pick = pd.DatetimeIndex(data['pickup_datetime'])
# creating an instance(date) of DatetimeIndex class using "dropoff_datetime"
date_drop = pd.DatetimeIndex(data['dropoff_datetime'])

# extracting new columns from "pick datetime"

# Last day of year when pickup was done
data['doy_pick'] = date_pick.dayofyear

# week of year when pickup was done
data['woy_pick'] = date_pick.weekofyear

# month of year when pickup was done
data['moy_pick'] = date_pick.month

# day of week when pickup was done
data['dow_pick'] = date_pick.dayofweek

# hour of day when pickup was done
data['hod_pick'] = date_pick.hour

# extracting new columns from "dropoff datetime"

# Last day of year dropoff was done
data['doy_drop'] = date_drop.dayofyear

# week of year when dropoff was done
data['woy_drop'] = date_drop.weekofyear

# month of year when dropoff was done
data['moy_drop'] = date_drop.month

# day of week when dropoff was done
data['dow_drop'] = date_drop.dayofweek

# hour of day when dropoff was done
data['hod_drop'] = date_drop.hour

C:\Users\vempa\AppData\Local\Temp\ipykernel_6968\1098005334.py:12: FutureWarning: weekofyear and week have been deprecated, please use DatetimeIndex.isocalendar().week
instead, which returns a Series. To exactly reproduce the behavior of week and weekofyear and return an Index, you may call pd.Int64Index(idx.isocalendar().week)
data['woy_pick'] = date_pick.weekofyear
C:\Users\vempa\AppData\Local\Temp\ipykernel_6968\1098005334.py:30: FutureWarning: weekofyear and week have been deprecated, please use DatetimeIndex.isocalendar().week
instead, which returns a Series. To exactly reproduce the behavior of week and weekofyear and return an Index, you may call pd.Int64Index(idx.isocalendar().week)
data['woy_drop'] = date_drop.weekofyear
```

```
In [21]: data.dtypes
```

```
Out[21]: id                object
vendor_id              int64
pickup_datetime        object
dropoff_datetime       object
passenger_count        int64
pickup_longitude       float64
pickup_latitude        float64
dropoff_longitude      float64
dropoff_latitude       float64
store_and_fwd_flag     object
trip_duration          int64
doy_pick               int64
woy_pick               int64
moy_pick               int64
dow_pick               int64
hod_pick               int64
doy_drop               int64
woy_drop               int64
moy_drop               int64
dow_drop               int64
hod_drop               int64
dtype: object
```

```
In [22]: data = pd.get_dummies(data.drop('id',axis=1), columns = ['store_and_fwd_flag'])
```

```
In [23]: data.tail()
```

Out[23]:

	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	trip_duration	doy_pick	...	moy_pick	dow_pick	hod_pick	doy_drop
729317	2	2016-05-21 13:29:38	2016-05-21 13:34:34	2	-73.965919	40.789780	-73.952637	40.789181	296	142	...	5	5	13	
729318	1	2016-02-22 00:43:11	2016-02-22 00:48:26	1	-73.996666	40.737434	-74.001320	40.731911	315	53	...	2	0	0	
729319	1	2016-04-15 18:56:48	2016-04-15 19:08:01	1	-73.997849	40.761696	-74.001488	40.741207	673	106	...	4	4	18	
729320	1	2016-06-19 09:50:47	2016-06-19 09:58:14	1	-74.006706	40.708244	-74.013550	40.713814	447	171	...	6	6	9	
729321	2	2016-01-01 17:24:16	2016-01-01 17:44:40	4	-74.003342	40.743839	-73.945847	40.712841	1224	1	...	1	4	17	

5 rows × 21 columns

```
In [25]: data_cleaned = data.drop(['pickup_datetime', 'dropoff_datetime'], axis=1)
```

Segregating variables: Independent and Dependent Variables

```
In [26]: #seperating independent and dependent variables
x = data.drop(['trip_duration','pickup_datetime','dropoff_datetime'], axis=1)
y = data['trip_duration']
x.shape, y.shape
```

Out[26]: ((693076, 18), (693076,))

Splitting the data into train set and the test set

```
In [27]: # Importing the train test split function
from sklearn.model_selection import train_test_split
train_x,test_x,train_y,test_y = train_test_split(x,y, random_state = 42)
```

Implementing Linear Regression

```
In [28]: #importing Linear Regression and metric mean square error
from sklearn.linear_model import LinearRegression as LR
from sklearn.metrics import mean_absolute_error as mae
from sklearn.metrics import r2_score
```

```
In [29]: # Creating instance of Linear Regression
lr = LR()

# Fitting the model
lr.fit(train_x, train_y)
```

```
Out[29]: ▾ LinearRegression
LinearRegression()
```

```
In [30]: # Predicting over the Train Set and calculating error
train_predict = lr.predict(train_x)
k = mae(train_predict, train_y)
print('Training Mean Absolute Error', k )
R_squared = r2_score(train_predict,train_y)
print('R2 score on train set', R_squared )
```

```
Training Mean Absolute Error 935.9382019789173
R2 score on train set 0.887845941706592
```

```
In [31]: # Predicting over the Test Set and calculating error
test_predict = lr.predict(test_x)
k = mae(test_predict, test_y)
print('Test Mean Absolute Error ', k )
R_squared = r2_score(test_predict,test_y)
print('R2 score on test set', R_squared )
```

```
Test Mean Absolute Error 933.0378564243354
R2 score on test set 0.8268991472379279
```

Parameters of Linear Regression


```
In [32]: lr.coef_
```

```
Out[32]: array([[ 1.97460507e+01,  1.00938024e+00,  8.25841764e+02, -3.60814511e+02,
  1.99492352e+02, -4.55845213e+02, -7.75964570e+04, -2.45443053e+01,
  1.87602423e+02,  2.12251139e+02, -3.22946315e+03,  7.75963920e+04,
  2.48700620e+01, -1.85726162e+02, -2.14279628e+02,  3.23019609e+03,
  3.81055761e+00, -3.81055761e+00])
```

Plotting the coefficients

```
In [33]: plt.figure(figsize=(8, 6), dpi=120, facecolor='w', edgecolor='b')
x = range(len(train_x.columns))
y = lr.coef_
plt.bar( x, y )
plt.xlabel( "Variables")
plt.ylabel('Coefficients')
plt.title('Coefficient plot')
```

```
Out[33]: Text(0.5, 1.0, 'Coefficient plot')
```



Here we can see that the model depends upon some Independent variables too much, But these coefficients are not suitable for interpretation because these are not scaled, therefore we will normalize and interpret later.

Checking assumptions of Linear Model

```
In [34]: # Arranging and calculating the Residuals
residuals = pd.DataFrame({
    'fitted values' : test_y,
    'predicted values' : test_predict,
})

residuals['residuals'] = residuals['fitted values'] - residuals['predicted values']
residuals.head()
```

```
Out[34]:
```

	fitted values	predicted values	residuals
22461	1550	95.523576	1454.476424
606882	2004	3411.637743	-1407.637743
509380	1611	94.518419	1516.481581
105494	91	96.257752	-5.257752
54890	730	3195.555017	-2465.555017

Plotting residual curve (Is there constant Variance OR Homoscedastic?)

```
In [35]: residuals.residuals[:]
```


```
Out[35]:
```

22461	1454.476424
606882	-1407.637743
509380	1516.481581
105494	-5.257752
54890	-2465.555017
...	
279147	188.122587
71621	-533.849650
242060	-2420.875546
152036	35.267649
665089	1229.155245

Name: residuals, Length: 173269, dtype: float64

```
In [36]: plt.figure(figsize=(10, 6), dpi=120, facecolor='w', edgecolor='b')
f = range(0,173269)
k = [0 for i in range(0,173269)]
plt.scatter( f, residuals.residuals[:], label = 'residuals')
plt.plot( f, k , color = 'red', label = 'regression line' )
plt.xlabel('fitted points ')
plt.ylabel('residuals')
plt.title('Residual plot')
plt.ylim(-6000, 6000)
plt.legend()
```


```
Out[36]: <matplotlib.legend.Legend at 0x15c897e1e80>
```



The Residual plot clearly Looks Homoscedastic, i.e. the the variance of the error across the dataset is nearly constant.

Checking Distribution of Residuals

```
In [37]: # Histogram for distribution
plt.figure(figsize=(10, 6), dpi=120, facecolor='w', edgecolor='b')
plt.hist(residuals.residuals, bins = 150)
plt.xlabel('Error')
plt.ylabel('Frequency')
plt.title('Distribution of Error Terms')
# plt.xlim(-10000, 20000)
plt.show()
```




According to the Histogram, the distribution of error is not normal

QQ-Plot (Is the data Normally Distributed?)

```
In [38]: # importing the QQ-plot from the from the statsmodels
from statsmodels.graphics.gofplots import qqplot

## Plotting the QQ plot
fig, ax = plt.subplots(figsize=(5,5) , dpi = 120)
qqplot(residuals.residuals, line = 's' , ax = ax)
plt.ylabel('Residual Quantiles')
plt.xlabel('Ideal Scaled Quantiles')
plt.title('Checking distribution of Residual Errors')
plt.show()
```

C:\Users\vempa\anaconda3\lib\site-packages\statsmodels\graphics\gofplots.py:993: UserWarning: marker is redundantly defined by the 'marker' keyword argument and the fm
t string "bo" (-> marker='o'). The keyword argument will take precedence.
ax.plot(x, y, fmt, **plot_style)



The QQ-plot clearly verifies our findings from the the histogram of the residuals, the data is not normal, Even after removing outliers with interquartile range method, there are still outliers all over the plot.

Variance Inflation Factor (VIF) (Checking for multi collinearity)

```
In [39]: # Importing Variance_inflation_Factor funtion from the Statsmodels
from statsmodels.stats.outliers_influence import variance_inflation_factor
from statsmodels.tools.tools import add_constant

# Calculating VIF for every column (only works for the not Catagorical)
VIF = pd.Series([variance_inflation_factor(data_cleaned.values, i) for i in range(data_cleaned.shape[1])], index =data_cleaned.columns)
VIF
```

```
Out[39]: vendor_id          1.097122e+00
passenger_count      1.090781e+00
pickup_longitude     1.100462e+00
pickup_latitude      1.203827e+00
dropoff_longitude     1.250740e+00
dropoff_latitude      1.380683e+00
trip_duration         8.978217e+00
doy_pick              1.047890e+07
woy_pick              6.959414e+02
moy_pick              9.132378e+03
dow_pick              9.525813e+01
hod_pick              2.811622e+02
doy_drop              1.047914e+07
woy_drop              6.972797e+02
moy_drop              9.132563e+03
dow_drop              9.535981e+01
hod_drop              2.880422e+02
store_and_fwd_flag_N  1.671513e+07
store_and_fwd_flag_Y  8.833919e+04
dtype: float64
```

```
In [40]: round(VIF[0:14],2)
```

```
Out[40]: vendor_id          1.10
passenger_count      1.09
pickup_longitude     1.10
pickup_latitude      1.20
dropoff_longitude     1.25
dropoff_latitude      1.38
trip_duration         8.98
doy_pick              10478897.96
woy_pick              695.94
moy_pick              9132.38
dow_pick              95.26
hod_pick              281.16
doy_drop              10479141.97
woy_drop              697.28
dtype: float64
```

From this list, we clearly see that there happens to be many Independent Variable over the value of 5, which means that there are Many features that exhibit the Multicollinearity in the dataset. Note that VIF only works for the Continuous Variables.

```
In [41]: data_cleaned.dtypes

Out[41]: vendor_id          int64
passenger_count          int64
pickup_longitude         float64
pickup_latitude          float64
dropoff_longitude        float64
dropoff_latitude         float64
trip_duration            int64
doy_pick                 int64
woy_pick                 int64
moy_pick                 int64
dow_pick                 int64
hod_pick                 int64
doy_drop                 int64
woy_drop                 int64
moy_drop                 int64
dow_drop                 int64
hod_drop                 int64
store_and_fwd_flag_N     uint8
store_and_fwd_flag_Y     uint8
dtype: object
```

Model Interpretability

So far we have simply been predicting the values using the linear regression, But in order to Interpret the model, the normalising of the data is essential.

```
In [42]: train_x['vendor_id']

Out[42]: 5342      2
326132    2
288391    1
205218    2
433541    1
..
272816    2
385069    2
138799    2
706259    1
128270    1
Name: vendor_id, Length: 519807, dtype: int64

In [43]: # Creating instance of Linear Regression
lr = LR(normalize = True)

# Fitting the model
lr.fit(train_x, train_y)

C:\Users\vempe\anaconda3\lib\site-packages\sklearn\linear_model\_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0 and will be removed in 1.2.
If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing stage. To reproduce the previous behavior:

from sklearn.pipeline import make_pipeline

model = make_pipeline(StandardScaler(with_mean=False), LinearRegression())

If you wish to pass a sample_weight parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

kwargs = {s[0] + '__sample_weight': sample_weight for s in model.steps}
model.fit(X, y, **kwargs)

warnings.warn(

Out[43]: LinearRegression
LinearRegression(normalize=True)
```

```
In [44]: # Predicting over the Train Set and calculating error
train_predict = lr.predict(train_x)
k = mae(train_predict, train_y)
print('Training Mean Absolute Error', k )
R_squared = r2_score(train_predict,train_y)
print('R2 score on test set', R_squared )

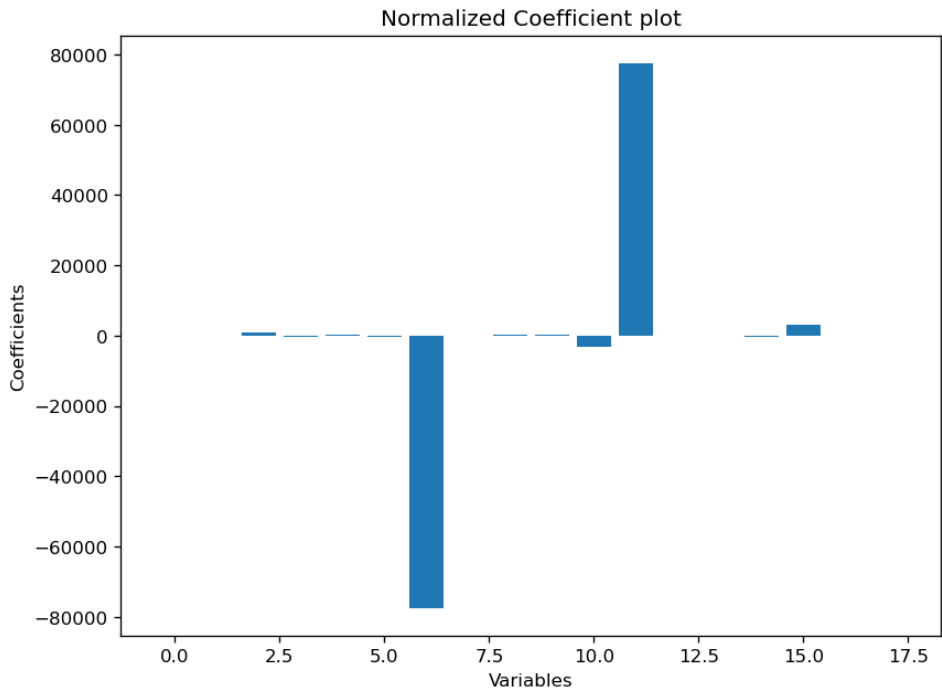
Training Mean Absolute Error 935.9382019791508
R2 score on test set 0.8878459417066005

In [45]: # Predicting over the Test Set and calculating error
test_predict = lr.predict(test_x)
k = mae(test_predict, test_y)
print('Test Mean Absolute Error ', k )
R_squared = r2_score(test_predict,test_y)
print('R2 score on test set', R_squared )

Test Mean Absolute Error 933.0378564245725
R2 score on test set 0.8268991472379501
```

```
In [46]: plt.figure(figsize=(8, 6), dpi=120, facecolor='w', edgecolor='b')
x = range(len(train_x.columns))
y = lr.coef_
plt.bar( x, y )
plt.xlabel( "Variables")
plt.ylabel('Coefficients')
plt.title('Normalized Coefficient plot')
```

Out[46]: Text(0.5, 1.0, 'Normalized Coefficient plot')



Now the coefficients we see are normalised and we can easily make final inferences out of it.

Here we can see that there are a lot of Coefficients which are near to zero and not Significant. So let us try removing them and build the model again.

Creating new subsets of data

```
In [47]: #seperating independent and dependent variables
x = data.drop(['trip_duration','pickup_datetime','dropoff_datetime'], axis=1)
y = data['trip_duration']
x.shape, y.shape
```

Out[47]: ((693076, 18), (693076,))

Arranging coefficients with features

```
In [48]: Coefficients = pd.DataFrame({
'Variable' : x.columns,
'coefficient' : lr.coef_
})
Coefficients.head()
```

Out[48]:

	Variable	coefficient
0	vendor_id	19.746051
1	passenger_count	1.009380
2	pickup_longitude	825.841764
3	pickup_latitude	-360.814511
4	dropoff_longitude	199.492352

Choosing variables with significance greater than 0.5 (Filtering Significant Features)

```
In [49]: sig_var = Coefficients[Coefficients.coefficient > 0.5]
```


Extracting the significant subset of independent Variables

```
In [50]: ▶ subset = data[sig_var['Variable'].values]
subset.head()
```

```
Out[50]:
```

	vendor_id	passenger_count	pickup_longitude	dropoff_longitude	moy_pick	dow_pick	doy_drop	woy_drop	hod_drop	store_and_fwd_flag_N
0	2	1	-73.953918	-73.963875	2	0	60	9	16	1
1	1	2	-73.988312	-73.994751	3	4	71	10	23	1
2	2	2	-73.997314	-73.948029	2	6	52	7	18	1
3	2	6	-73.961670	-73.956779	1	1	5	1	10	1
4	1	1	-74.017120	-73.988182	2	2	48	7	6	1

Splitting the data into train set and the test set

```
In [51]: ▶ # Importing the train test split function
from sklearn.model_selection import train_test_split
train_x, test_x, train_y, test_y = train_test_split(subset, y, random_state = 42)
```

Implementing Linear Regression

```
In [52]: ▶ #importing Linear Regression and metric mean square error
from sklearn.linear_model import LinearRegression as LR
from sklearn.metrics import mean_absolute_error as mae
```

Training Model

```
In [53]: ▶ # Creating instance of Linear Regression with Normalised Data
lr = LR(normalize = True)

# Fitting the model
lr.fit(train_x, train_y)
```

C:\Users\vempa\anaconda3\lib\site-packages\sklearn\linear_model_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0 and will be removed in 1.2. If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing stage. To reproduce the previous behavior:

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False), LinearRegression())
```

If you wish to pass a sample_weight parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '_sample_weight': sample_weight for s in model.steps}
model.fit(X, y, **kwargs)
```

```
Out[53]:
```

	LinearRegression
LinearRegression(normalize=True)	

Predicting over the train set

```
In [54]: ▶ # Predicting over the Train Set and calculating error
train_predict = lr.predict(train_x)
k = mae(train_predict, train_y)
print('Training Mean Absolute Error', k )
```

Training Mean Absolute Error 560.9403158158182

Predicting over the test set


```
In [55]: ▶ # Predicting over the Test Set and calculating error
test_predict = lr.predict(test_x)
k = mae(test_predict, test_y)
print('Test Mean Absolute Error', k )
```

Test Mean Absolute Error 556.8090417448929

Plotting the coefficients

```
In [56]: plt.figure(figsize=(8, 6), dpi=120, facecolor='w', edgecolor='b')
x = range(len(train_x.columns))
y = lr.coef_
plt.bar( x, y )
plt.xlabel( "Variables")
plt.ylabel('Coefficients')
plt.title('Normalized Coefficient plot')
```

Out[56]: Text(0.5, 1.0, 'Normalized Coefficient plot')



Ridge Regression

```

In [66]: from sklearn.linear_model import Ridge
from sklearn.metrics import r2_score

ridgeRegressor = Ridge(alpha = 0.5 )
ridgeRegressor.fit(train_x,train_y)
y_test_predict_ridge = ridgeRegressor.predict(test_x)
y_train_predict_ridge = ridgeRegressor.predict(train_x)

# R_squared = r2_score(y_predict_ridge,test_y)
# print('R2 Score on test set', R_squared )

# Predicting over the Train Set and calculating error
# train_predict = lr.predict(train_x)
k1 = mae(y_train_predict_ridge, train_y)
print('Training Mean Absolute Error', k1 )

#### Predicting over the test set

# Predicting over the Test Set and calculating error
# test_predict = lr.predict(test_x)
k2 = mae(y_test_predict_ridge, test_y)
print('Test Mean Absolute Error ', k2 )

coefficient_df = pd.DataFrame()
coefficient_df['Column_Name'] = train_x.columns
coefficient_df['Coefficient_Value'] = pd.Series(ridgeRegressor.coef_)
print(coefficient_df.head(15))

plt.rcParams["figure.figsize"] = (15,6)
plt.bar(coefficient_df['Column_Name'],coefficient_df['Coefficient_Value'])

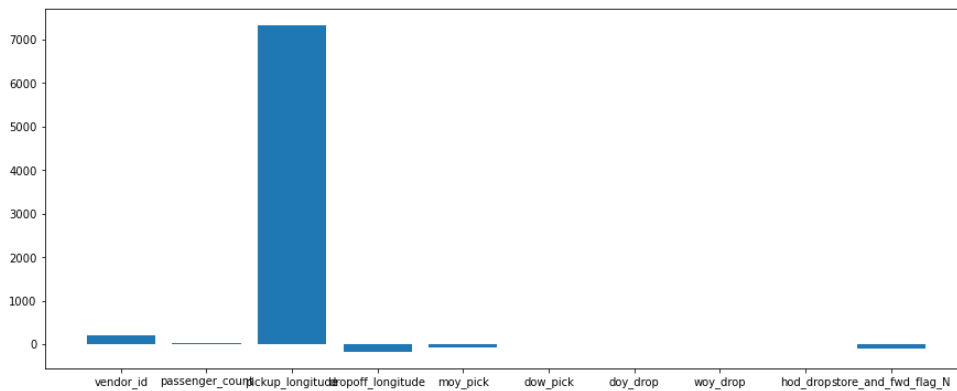
```

```

Training Mean Absolute Error 560.9328633555573
Test Mean Absolute Error    556.8012764347889
   Column_Name  Coefficient_Value
0      vendor_id          189.728970
1  passenger_count          16.085504
2  pickup_longitude      7335.928229
3  dropoff_longitude     -179.946308
4         moy_pick        -71.467552
5         dow_pick         0.076024
6         doy_drop         2.945474
7         woy_drop         0.056995
8         hod_drop         1.224168
9  store_and_fwd_flag_N     -99.704619

```

Out[66]: <BarContainer object of 10 artists>



```

In [67]: from sklearn.linear_model import Ridge
from sklearn.metrics import r2_score

list2 = [0.5,1,5,10,50,100]

for i in list2:
    ridgeRegressor = Ridge(alpha = i )
    ridgeRegressor.fit(train_x,train_y)
    # y_predict_ridge = ridgeRegressor.predict(test_x)
    y_test_predict_ridge = ridgeRegressor.predict(test_x)
    y_train_predict_ridge = ridgeRegressor.predict(train_x)

    # R_squared = r2_score(y_predict_ridge,test_y)
    # print('R2 Score on test set', R_squared )

    # Predicting over the Train Set and calculating error
    # train_predict = lr.predict(train_x)
    k1 = mae(y_train_predict_ridge, train_y)
    print('Training Mean Absolute Error', k1 )

    #### Predicting over the test set

    # Predicting over the Test Set and calculating error
    # test_predict = lr.predict(test_x)
    k2 = mae(y_test_predict_ridge, test_y)
    print('Test Mean Absolute Error ', k2 )

    coefficient_df = pd.DataFrame()
    coefficient_df['Column_Name'] = train_x.columns
    coefficient_df['Coefficient_Value'] = pd.Series(ridgeRegressor.coef_)
    print(coefficient_df.head(15))

    plt.rcParams["figure.figsize"] = (15,6)
    plt.bar(coefficient_df['Column_Name'],coefficient_df['Coefficient_Value'])

```

Training Mean Absolute Error 560.932863355573

Test Mean Absolute Error 556.8012764347889

	Column_Name	Coefficient_Value
0	vendor_id	189.728970
1	passenger_count	16.085504
2	pickup_longitude	7335.928229
3	dropoff_longitude	-179.946308
4	moy_pick	-71.467552
5	dow_pick	0.076024
6	doy_drop	2.945474
7	woy_drop	0.056995
8	hod_drop	1.224168
9	store_and_fwd_flag_N	-99.704619

Training Mean Absolute Error 560.9254744431873

Test Mean Absolute Error 556.7935956156483

	Column_Name	Coefficient_Value
0	vendor_id	189.735586
1	passenger_count	16.085377
2	pickup_longitude	7329.996764
3	dropoff_longitude	-177.754599
4	moy_pick	-71.466526
5	dow_pick	0.073072
6	doy_drop	2.945426
7	woy_drop	0.057266
8	hod_drop	1.225140
9	store_and_fwd_flag_N	-99.748400

Training Mean Absolute Error 560.8684072619743

Test Mean Absolute Error 556.734181647583

	Column_Name	Coefficient_Value
0	vendor_id	189.788113
1	passenger_count	16.084358
2	pickup_longitude	7282.927011
3	dropoff_longitude	-160.521745
4	moy_pick	-71.458304
5	dow_pick	0.049619
6	doy_drop	2.945038
7	woy_drop	0.059420
8	hod_drop	1.232827
9	store_and_fwd_flag_N	-100.094594

Training Mean Absolute Error 560.8016783831772

Test Mean Absolute Error 556.6649950914696

	Column_Name	Coefficient_Value
0	vendor_id	189.852770
1	passenger_count	16.083074
2	pickup_longitude	7225.026697
3	dropoff_longitude	-139.711222
4	moy_pick	-71.447980
5	dow_pick	0.020696
6	doy_drop	2.944554
7	woy_drop	0.062077
8	hod_drop	1.242203
9	store_and_fwd_flag_N	-100.517385

Training Mean Absolute Error 560.4343604536364


Test Mean Absolute Error 556.2840884199773

	Column_Name	Coefficient_Value
0	vendor_id	190.332841
1	passenger_count	16.072598
2	pickup_longitude	6795.948003
3	dropoff_longitude	1.434118
4	moy_pick	-71.363948
5	dow_pick	-0.196121
6	doy_drop	2.940712
7	woy_drop	0.082056
8	hod_drop	1.308996
9	store_and_fwd_flag_N	-103.534618

Training Mean Absolute Error 560.2968631920423

Test Mean Absolute Error 556.1405541289549

	Column_Name	Coefficient_Value
0	vendor_id	190.851843
1	passenger_count	16.059593
2	pickup_longitude	6331.565163
3	dropoff_longitude	129.148546
4	moy_pick	-71.256822
5	dow_pick	-0.435524
6	doy_drop	2.936015
7	woy_drop	0.104204
8	hod_drop	1.376135
9	store_and_fwd_flag_N	-106.528160



MAE reduced a bit for Ridge regressor

As we increase the alpha, the coefficients are changing but not becoming zero completely as expected in lasso model

Bar plot shows that pick longitude, drop off longitude has significant impact on the target variable - trip duration

surprisingly pick up latitude, dropoff latitude has no significant impact.

storeand flag co-efficient is showing impact on duration which is not practical

pickup longitude is more significant than the date and time - maybe the location(Traffic,roads) could be the major factor

Lasso Regression

```
In [64]: from sklearn.linear_model import Lasso
from sklearn.metrics import r2_score

LassoRegressor = Lasso(alpha = 0.5)
LassoRegressor.fit(train_x, train_y)
y_test_predict_lasso = LassoRegressor.predict(test_x)
y_train_predict_lasso = LassoRegressor.predict(train_x)

k3 = mae(y_train_predict_lasso, train_y)
print('Training Mean Absolute Error', k1 )


k4 = mae(y_test_predict_lasso, test_y)
print('Test Mean Absolute Error ', k2 )

coefficient_df = pd.DataFrame()
coefficient_df['Column_Name'] = train_x.columns
coefficient_df['Coefficient_Value'] = pd.Series(LassoRegressor.coef_)
print(coefficient_df.head(15))

plt.rcParams["figure.figsize"] = (15,6)
plt.bar(coefficient_df['Column_Name'], coefficient_df['Coefficient_Value'], alpha = 0.5)
```

```
Training Mean Absolute Error 580.6874771869304
Test Mean Absolute Error    576.2106723643577
   Column_Name  Coefficient_Value
0   vendor_id         187.203868
1  passenger_count         16.000716
2  pickup_longitude        6920.123193
3  dropoff_longitude         -0.000000
4     moy_pick          -64.780914
5     dow_pick          -0.000000
6     doy_drop           2.737195
7     woy_drop           0.010440
8     hod_drop           1.291366
9  store_and_fwd_flag_N         -8.556934
```

Out[64]: <BarContainer object of 10 artists>



```

In [65]: from sklearn.linear_model import Lasso
from sklearn.metrics import r2_score

list1 = [0.5,1,5,10,50,100]

for i in list1:

    LassoRegressor = Lasso(alpha = i )
    LassoRegressor.fit(train_x,train_y)
    y_test_predict_lasso = LassoRegressor.predict(test_x)
    y_train_predict_lasso = LassoRegressor.predict(train_x)

    k3 = mae(y_train_predict_lasso, train_y)
    print('Training Mean Absolute Error', k1 )

    k4 = mae(y_test_predict_lasso, test_y)
    print('Test Mean Absolute Error ', k2 )

    coefficient_df = pd.DataFrame()
    coefficient_df['Column_Name'] = train_x.columns
    coefficient_df['Coefficient_Value'] = pd.Series(LassoRegressor.coef_)
    print(coefficient_df.head(15))

    plt.rcParams["figure.figsize"] = (15,6)
    plt.bar(coefficient_df['Column_Name'],coefficient_df['Coefficient_Value'],alpha = 0.5)

```

Training Mean Absolute Error 580.6874771869304

Test Mean Absolute Error 576.2106723643577

	Column_Name	Coefficient_Value
0	vendor_id	187.203868
1	passenger_count	16.000716
2	pickup_longitude	6920.123193
3	dropoff_longitude	-0.000000
4	moy_pick	-64.780914
5	dow_pick	-0.000000
6	doy_drop	2.737195
7	woy_drop	0.010440
8	hod_drop	1.291366
9	store_and_fwd_flag_N	-8.556934

Training Mean Absolute Error 580.6874771869304

Test Mean Absolute Error 576.2106723643577

	Column_Name	Coefficient_Value
0	vendor_id	185.690914
1	passenger_count	15.889112
2	pickup_longitude	6535.353194
3	dropoff_longitude	0.000000
4	moy_pick	-58.124709
5	dow_pick	-0.050373
6	doy_drop	2.525743
7	woy_drop	-0.000000
8	hod_drop	1.317949
9	store_and_fwd_flag_N	-0.000000

Training Mean Absolute Error 580.6874771869304

Test Mean Absolute Error 576.2106723643577

	Column_Name	Coefficient_Value
0	vendor_id	174.379631
1	passenger_count	14.999367
2	pickup_longitude	3453.626906
3	dropoff_longitude	0.000000
4	moy_pick	-4.903665
5	dow_pick	-0.639785
6	doy_drop	0.836382
7	woy_drop	-0.095635
8	hod_drop	1.524114
9	store_and_fwd_flag_N	-0.000000

Training Mean Absolute Error 580.6874771869304

Test Mean Absolute Error 576.2106723643577

	Column_Name	Coefficient_Value
0	vendor_id	159.634495
1	passenger_count	13.909178
2	pickup_longitude	0.000000
3	dropoff_longitude	0.000000
4	moy_pick	-0.000000
5	dow_pick	-1.268891
6	doy_drop	0.677285
7	woy_drop	-0.000000
8	hod_drop	1.753835
9	store_and_fwd_flag_N	-0.000000

Training Mean Absolute Error 580.6874771869304

Test Mean Absolute Error 576.2106723643577

	Column_Name	Coefficient_Value
0	vendor_id	3.675598
1	passenger_count	7.626249
2	pickup_longitude	0.000000
3	dropoff_longitude	0.000000
4	moy_pick	-0.000000
5	dow_pick	-0.000000
6	doy_drop	0.655778
7	woy_drop	0.000000
8	hod_drop	0.943642
9	store_and_fwd_flag_N	-0.000000

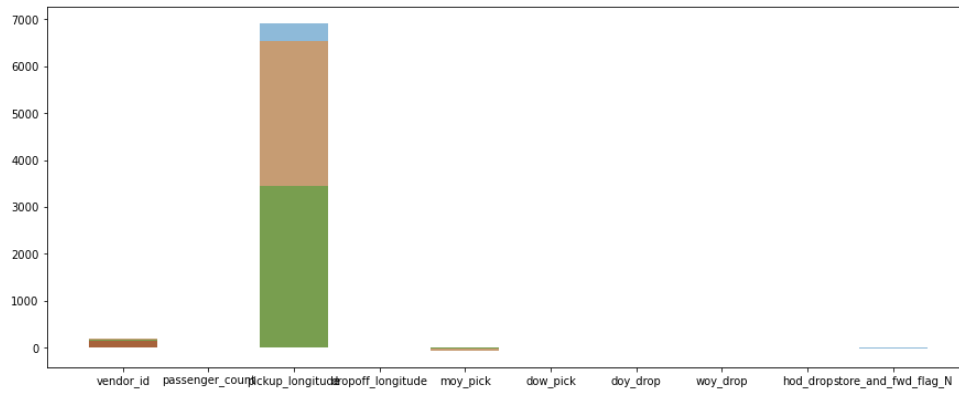
Training Mean Absolute Error 580.6874771869304

Test Mean Absolute Error 576.2106723643577

	Column_Name	Coefficient_Value
0	vendor_id	0.000000
1	passenger_count	0.000000
2	pickup_longitude	0.000000
3	dropoff_longitude	0.000000
4	moy_pick	-0.000000
5	dow_pick	-0.000000
6	doy_drop	0.636098
7	woy_drop	0.000000
8	hod_drop	0.000000
9	store_and_fwd_flag_N	-0.000000

Training Mean Absolute Error 580.6874771869304

Test Mean Absolute Error 576.2106723643577



As we increase the alpha, the coefficients are becoming zero as expected in lasso model

Bar plot shows that pick longitude has significant impact on the target variable - trip duration

surprisingly pick up latitude has no significant impact

pickup longitude is more significant than the date and time - maybe the location(Traffic,roads) could be the major factor

In []:

In []: